

Lab 5 report

In-Lab Tasks:

In lab Task 1:

```
# Download the data using wget
#!wget
"https://www.dropbox.com/s/veak3ugc4wj9luz/Alumni%20Giving%20Regression%20%28Edited%209.csv"

from keras.models import Sequential
from keras.layers import Dense, Dropout
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import numpy as np
from sklearn import linear_model
from sklearn import preprocessing
from sklearn import tree
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
import pandas as pd
import csv
import matplotlib.pyplot as plt
```

In lab Task 2:

```
np.random.seed(7)
df = pd.read_csv("Alumni Giving Regression.csv", delimiter=",")
dd_df_1 = df.head()
```

In lab Task 3:

```
# Describing the data
data_description = df.describe()
print("Data Description:\n", data_description)
```

Data Description:

	A	B	C	D	E	F
count	123.000000	123.000000	123.000000	123.000000	123.000000	123.000000
mean	17.772358	0.403659	0.136260	0.645203	0.841138	0.141789
std	4.517385	0.133897	0.060101	0.169794	0.083942	0.080674
min	6.000000	0.140000	0.000000	0.260000	0.580000	0.020000
25%	16.000000	0.320000	0.095000	0.505000	0.780000	0.080000
50%	18.000000	0.380000	0.130000	0.640000	0.840000	0.130000
75%	20.000000	0.460000	0.180000	0.785000	0.910000	0.170000
max	31.000000	0.950000	0.310000	0.960000	0.980000	0.410000

In lab Task 4:

```
#Compute Correlation

corr=df.corr(method = 'pearson')
corr
print(corr)
```

	A	B	C	D	E	F
A	1.000000	-0.691900	0.414978	-0.604574	-0.521985	-0.549244
B	-0.691900	1.000000	-0.581516	0.487248	0.376735	0.540427
C	0.414978	-0.581516	1.000000	0.017023	0.055766	-0.175102
D	-0.604574	0.487248	0.017023	1.000000	0.934396	0.681660
E	-0.521985	0.376735	0.055766	0.934396	1.000000	0.647625
F	-0.549244	0.540427	-0.175102	0.681660	0.647625	1.000000

In lab Task 5:

```
#Splitting Datasets
# Y_POSITION is set to 5, indicating that the target variable is in the
5th column
Y_POSITION = 5

# model_1_features is created as a list containing the indices [0, 1, 2,
3, 4],
# representing the columns from 0 to 4 (excluding 5).
model_1_features = [i for i in range(0, Y_POSITION)]

# Extract features (X) and target variable (Y) from the DataFrame (df)
X = df.iloc[:, model_1_features]
Y = df.iloc[:, Y_POSITION]
```

```
# Create model
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.20, random_state=2020)
```

In lab Task 6:

```
#Linear Regression
from sklearn.linear_model import LinearRegression

# Assuming df is our dataframe

# Y_POSITION is set to 5, indicating that the target variable is in the
5th column of the DataFrame.
Y_POSITION = 5

# model_1_features is created as a list containing the indices [0, 1, 2,
3, 4], representing the columns from 0 to 4 (excluding 5).
model_1_features = [i for i in range(0, Y_POSITION)]

# Extract features (X) and target variable (Y) from the DataFrame (df)
X = df.iloc[:, model_1_features]
Y = df.iloc[:, Y_POSITION]

# Create train/test split,given that we took testing dataset as 20% and
radnom state to be 2020 to generate reproduceability
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.20, random_state=2020)

# Model 1: Linear Regression
modell = LinearRegression()

# Train the model
modell.fit(X_train, y_train)

# Predictions on the training set
y_pred_train1 = modell.predict(X_train)
RMSE_train1 = mean_squared_error(y_train, y_pred_train1)

# Predictions on the testing set
y_pred1 = modell.predict(X_test)
```

```

RMSE_test1 = mean_squared_error(y_test, y_pred1)

# Display RMSE for both training and testing sets
print("Regression Train set: RMSE {}".format(RMSE_train1))
print("Regression Test set: RMSE {}".format(RMSE_test1))

# Display coefficients
coef_dict = {}
for coef, feat in zip(model1.coef_, model_1_features):
    coef_dict[df.columns[feat]] = coef
print("Coefficients:", coef_dict)


# Scatter plot of actual vs predicted values
x_values = np.arange(len(y_test))
plt.scatter(x_values, y_test, color='red', label='Actual')
plt.scatter(x_values, y_pred1, color='green', label='Predicted')
plt.xlabel('Index or Sequence of Values')
plt.ylabel('Values')
plt.title('Actual vs Predicted Values')
plt.legend()
plt.show()

```

Regression Train set: RMSE 0.002761693322289229

Regression Test set: RMSE 0.004209824026356377

Coefficients: {'A': -0.0009337757382416938, 'B': 0.16012156890162943, 'C': -0.044160015425349614, 'D': 0.15217907817100407, 'E': 0.17539950794101047}



Post lab :

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score

# Load the dataset
data = pd.read_csv("Salary_dataset.csv")

# Extracting features and target variable
X = data.iloc[:, :-1].values # Assuming the independent variable is in
the first column
y = data.iloc[:, -1].values # Assuming the dependent variable (salary)
is in the last column

# Splitting the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Linear Regression
model = LinearRegression()
model.fit(X_train, y_train)

# Model Evaluation
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Linear Regression Test MSE: {mse}")

# Polynomial Regression (Degree 2)
poly = PolynomialFeatures(degree=2)
X_poly = poly.fit_transform(X_train)
model_poly = LinearRegression()
model_poly.fit(X_poly, y_train)

# Model Evaluation - Polynomial Regression
X_test_poly = poly.transform(X_test)
y_pred_poly = model_poly.predict(X_test_poly)
mse_poly = mean_squared_error(y_test, y_pred_poly)
print(f"Polynomial Regression (Degree 2) Test MSE: {mse_poly}")

# Cross-validation
```

```
cross_val_scores = cross_val_score(model, X, y, cv=5,
scoring='neg_mean_squared_error')
print(f"Cross-validated MSE for Linear Regression: {-
cross_val_scores.mean()}")

# Cross-validation for Polynomial Regression
cross_val_scores_poly = cross_val_score(model_poly, poly.transform(X),
y, cv=5, scoring='neg_mean_squared_error')
print(f"Cross-validated MSE for Polynomial Regression: {-
cross_val_scores_poly.mean()}")
```

- **Salary_dataset.csv**(text/csv) - 664 bytes, last modified: 16/11/2023 - 100% done
Saving Salary_dataset.csv to Salary_dataset (3).csv
Linear Regression Test MSE: 55494098.13142202
Polynomial Regression (Degree 2) Test MSE: 64168200.124699205
Cross-validated MSE for Linear Regression: 78625892.88185735
Cross-validated MSE for Polynomial Regression: 70662205.86009356